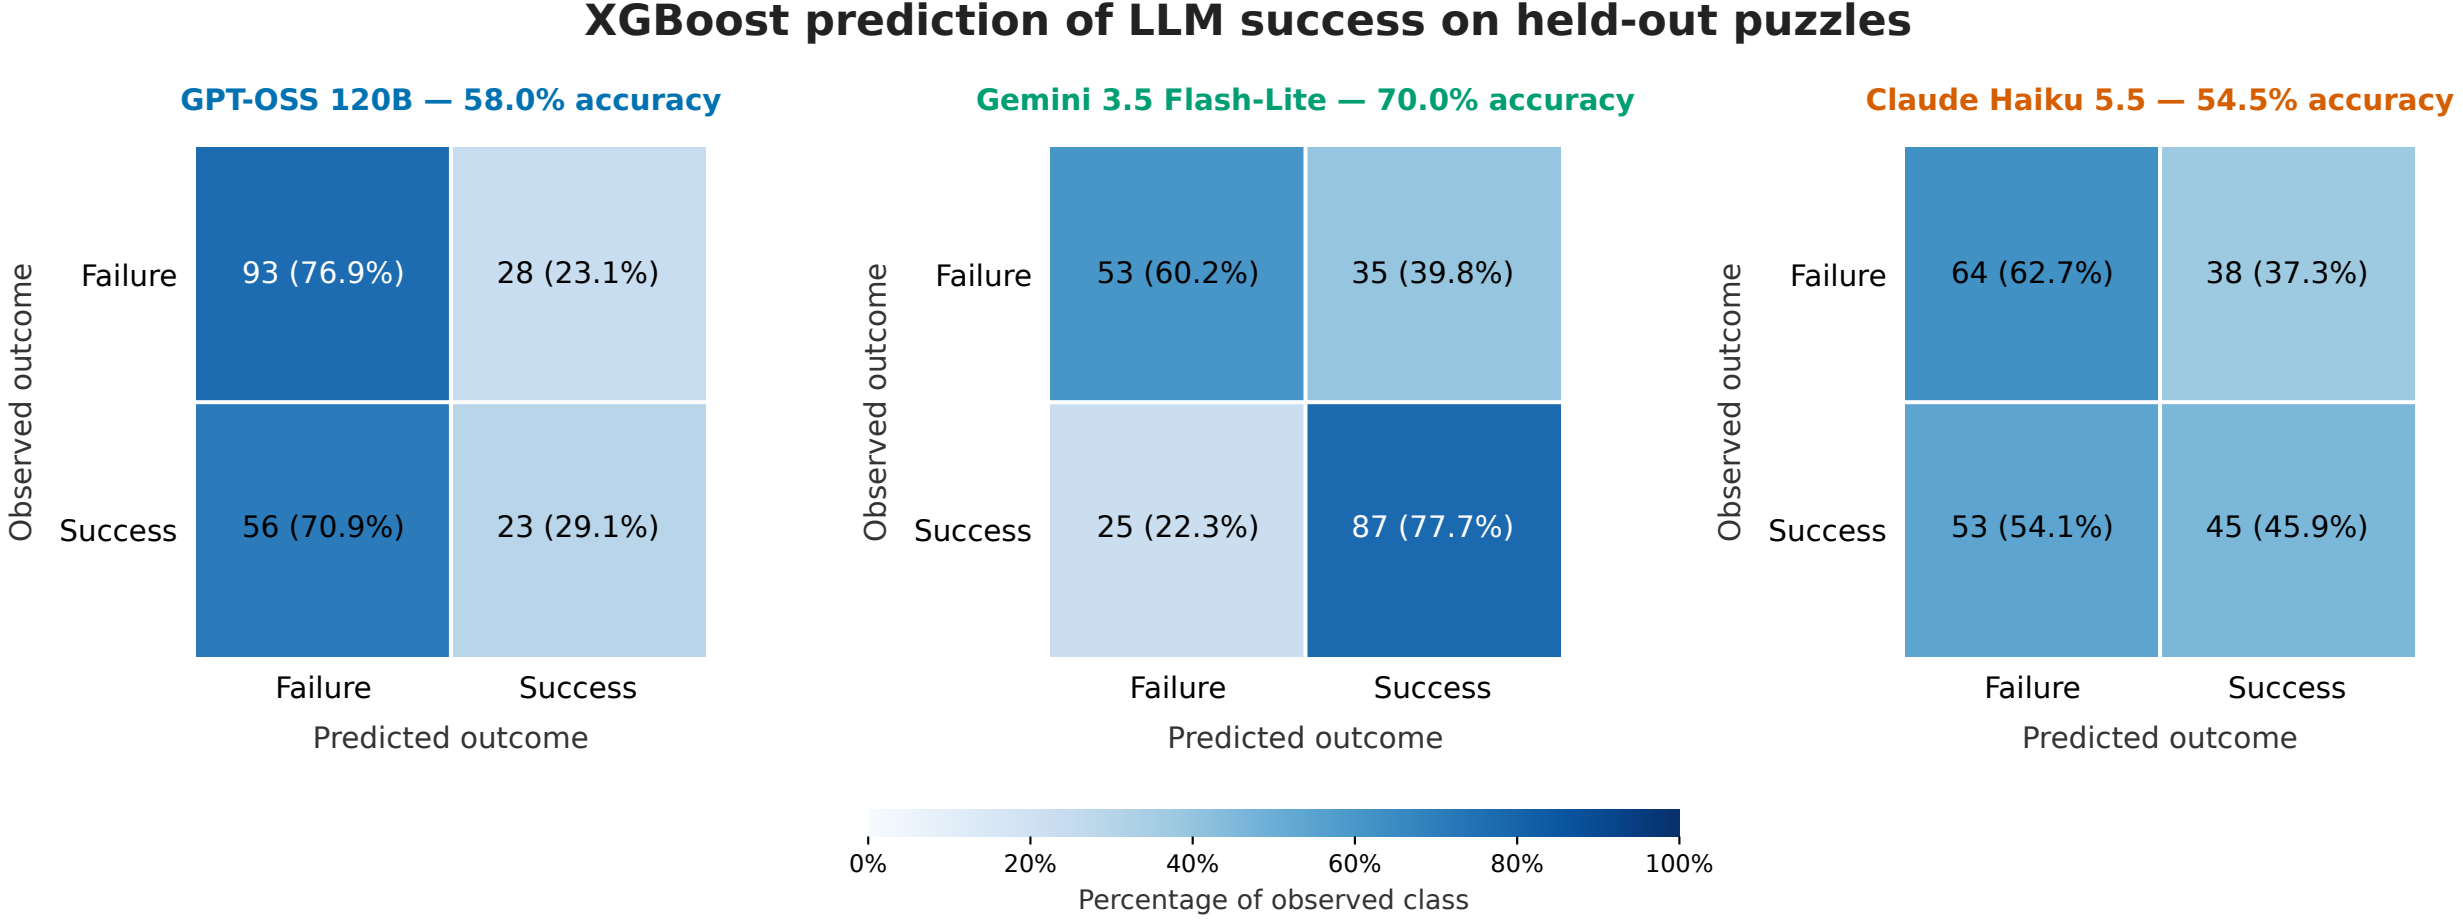
XGBoost prediction of LLM success on held-out puzzles
GPT-OSS 120B — 58.0% accuracy
Gemini 3.5 Flash-Lite — 70.0% accuracy
Claude Haiku 5.5 — 54.5% accuracy
Observed outcome
Failure
Success
93 (76.9%)
28 (23.1%)
56 (70.9%)
23 (29.1%)
53 (60.2%)
35 (39.8%)
25 (22.3%)
87 (77.7%)
64 (62.7%)
38 (37.3%)
53 (54.1%)
45 (45.9%)
Predicted outcome
0%
20%
40%
60%
80%
100%
Percentage of observed class